Understanding AI Cases in the U.S. Courts

Dawson Petersen **Eilat Herman** **Katherine Conhaim** **Kevin Tobia** **Nathan Schneider**

Georgetown University

1 Introduction

The number of U.S. court cases involving AI has increased exponentially in recent years. The current study investigates 1) what types of AI-related issues are arising in U.S. courts and 2) how the courts are construing AI systems. We find that although much academic discourse has focused on public-facing AI harms (such as deepfakes, copyright abuse, and algorithmic discrimination), courts are primarily focused on corporate disputes and AI assisted litigation. Additionally, we find that although laypeople frequently anthropomorphize AI systems (Mitchell and Krakauer, 2023), construing them as responsible (Petersen and Almor, 2025) animate agents, the courts draw clear linguistic distinctions between entities with legal personhood (people, corporations) and those without (AIs).

2 Methods

We searched Westlaw for the phrase "artificial intelligence" in opinions from all U.S. jurisdictions and levels. The search returned 1,959 cases from December 1988 to March 2026. From these opinions, we selected sentences which contained AI-related search terms (e.g., LLM, deep learning, chatbot, etc.), resulting in a corpus of 9,262 AI mentions from 1,938 cases (4.78 per case).

2.1 Topic Classification

To understand the relevance of AI to these cases, we established an inventory of topic categories and developed a pipeline for automatic classification using a weighted ensemble of LoRA-finetuned open source LLMs. First, we produced a training dataset by manually classifying 600 items into seven categories (*AI Assisted Litigation*; *AI Harm*; *AI, Society, and Law*; *Copyright Dispute*; *Data Scraping and Privacy*; *Trade Dispute*; and *Other*). Two independent raters achieved an agreement of 69% ($\kappa = .64$). Disagreements were adjudicated by a third rater. We used this dataset to finetune LoRA adapters for seven open source LLMs, using an 85/15 training-evaluation split. After inference, we combined models' classifications of the full dataset using F1-weighted ensemble voting. Items with low inter-model agreement (n = 673) were re-categorized as *Other*. To evaluate model performance, we created a test set by manually classified another 105 mentions. The ensemble model achieved an accuracy of 82% ($\kappa = .79$) on the test set.

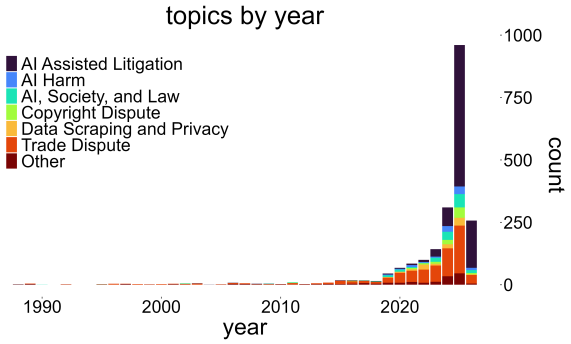


Figure 1: Case topics by year. Data for 2026 only includes cases which included in Westlaw by March 13th, 2026.

2.2 Anthropomorphism

To understand how the courts construe AI systems, we seperately analyzed the corpus using AnthroScore (Cheng et al., 2024), a RoBERTa-based measure of whether an entity is being described as an animate agent or an object. We computed AnthroScores for phrases containing terms denoting AI (AI, LLM, intelligence) as well as four reference categories: humans (person, individual), companies (company, corporation), documents (policy, contract), and problems (problem, issue).

Entity Type	Mean AnthroScore
Human	−0.12
Company	−0.66
AI	−2.15
Document	−2.38
Problem	−1.97

Table 1: Mean AnthroScores by entity type. Lower numbers indicate less anthropomorphic descriptions.

3 Results

Prior to 2023, most AI mentions (58%) were *Trade Disputes*. While the number of *Trade Dispute* mentions continues to grow, they make up minority of new mentions (27%), with the majority instead being *AI Assisted Litigation* (57%) (see Figure 1). We also found that judicial opinions show a clear linguistic contrast between entities with legal personhood and those without, with AnthroScores for AI terms patterning more closely with other inanimate entities than with humans or companies (see Table 1). Our findings underscore the widespread impact of AI on the legal system as a legal tool and a subject of corporate litigation, although the most public facing AI harms have had a more limited impact. Our findings further allay some concerns regarding AI anthropomorphism, by showing that it is not highly prevalent in judicial opinions.

References

Myra Cheng, Kristina Gligorić, Tiziano Piccardi, and Dan Jurafsky. 2024. Anthroscore: A computational linguistic measure of anthropomorphism. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 807–825.

Melanie Mitchell and David C Krakauer. 2023. The debate over understanding in ai's large language models. *Proceedings of the National Academy of Sciences*, 120(13):e2215907120.

Dawson Petersen and Amit Almor. 2025. Agentive linguistic framing affects responsibility assignments toward ais and their creators. *Frontiers in Psychology*, 16:1498958.